

DASC5301 Data Science, Fall 2021, Chengkai Li, University of Texas at Arlington

Programming Assignment 3 Solution

Due: Friday, December 3, 2021, 11:59pm

Academic Honesty

1. This assignment must be done individually and independently. You must implement the whole assignment by yourself. Academic dishonesty is not tolerated.
2. You can discuss topics related to the assignment with your fellow students. But you are not allowed to discuss/share your solution and code.

Requirements

1. When you work on this assignment, you should make a copy of this notebook in Google Colab. This can be done using the option File > Save a copy in Drive in Google Colab.
2. You should fill in your answer for each task inside the code block right below the task.
3. You should only insert your code into the designated code blocks, as mentioned above. Other than that, you shouldn't change anything else in the notebook, unless otherwise instructed.
4. For each code block, you are free to use multiple lines of code. Tasks 1-4 only need one line (or a few) each. Tasks 5 and 6 will need longer codes.
5. Even if you can only partially solve a task, you should include your code in the code block, which allows us to consider partial credit.
6. However, your code should not raise errors. Any code raising errors will not get partial credit.
7. We will test your code in Google Colab. Make sure your code runs in Google Colab.
8. For classification, you are expected to use `scikit-learn` (http://scikit-learn.org/stable/supervised_learning.html#supervised-learning). Refer to our Colab on classification for a tutorial of how to use `scikit-learn` to build classification models. The Colab has been on the Syllabus page and has been explained during lectures. Its link is https://colab.research.google.com/drive/1_1N7Hz3-mM2GAatME1JBEMW5ibqJpewe.

9. For feature extraction and data manipulation, you are expected to use `pandas` and `NumPy`, which we studied in the first half of this semester.
10. To submit your assignment, download your Colab into a `.ipynb` file. This can be done using the option `Download > Download .ipynb` in Google Colab.
11. Submit the downloaded `.ipynb` file into the Programming Assignment 3 entry in Canvas.

▼ Datasets

In this assignment, we will use Python `scikit-learn` to build a classifier on a dataset about data science workers and learners. This dataset has close to three thousand rows and quite many complex columns. To make it easier to get started, we also provide a smaller dataset with less columns. Both datasets are provided as CSV files in the assignment's entry in canvas. You will need to upload these CSV files to your Google Colab working directory. Once the CSV files are in your working directory, let's load the small CSV file `small_ds_workers_learners.csv` into a `pandas DataFrame`.

```
import pandas as pd
survey = pd.read_csv('small_ds_workers_learners.csv', delimiter=',', decimal=",")
```

Let's gain some basic understanding of the dataset by using `info()`.

```
survey.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2650 entries, 0 to 2649
Data columns (total 4 columns):
 #   Column                                                                 Non-Null Count  Dtype
---  -
 0   Yearly salary                                                         1845 non-null   object
 1   Years of experience in machine learning methods                    2462 non-null   object
 2   Most frequently used big data products                             624 non-null    object
 3   Regularly use Scikit-learn                                          1433 non-null   object
dtypes: object(4)
memory usage: 82.9+ KB
```

We can see that there are null values in every column, since the non-null count of each column is less than 2650. We will make `Yearly salary` the class/prediction attribute. Therefore, let's go ahead to remove rows with missing values in column `Yearly salary`.

```
survey = survey[survey['Yearly salary'].notna()]
```

Now let's find out all distinct values in column Yearly salary.

```
survey['Yearly salary'].unique()
```

```
array(['15,000-19,999', '100,000-124,999', '70,000-79,999',  
      '300,000-499,999', '200,000-249,999', '125,000-149,999',  
      '60,000-69,999', '25,000-29,999', '250,000-299,999',  
      '80,000-89,999', '40,000-49,999', '150,000-199,999', '0-999',  
      '90,000-99,999', '30,000-39,999', '50,000-59,999', '10,000-14,999',  
      '500,000-999,999', '4,000-4,999', '1,000-1,999', '2,000-2,999',  
      '>1,000,000', '20,000-24,999', '7,500-9,999', '5,000-7,499',  
      '3,000-3,999'], dtype=object)
```

The classification task in this assignment is to predict whether a data science worker/learner makes more than \$100K in a year or not, i.e., it is a binary classification task. Hence, we now replace all salary values less than \$100K with 'No', and replace all other values with 'Yes'.

```
survey['Yearly salary'] = survey['Yearly salary'].map({'100,000-149,999': 'Yes', '150,000-199,999': 'Yes',  
survey.loc[survey['Yearly salary'] != 'Yes', 'Yearly salary'] = 'No'
```

Now let's take a look at the first 20 rows after these transformations.

```
survey.head(20)
```

	Yearly salary	Years of experience in machine learning methods	Most frequently used big data products	Regularly use Scikit-learn
1	No	3-4 years	MySQL	Yes
3	No	4-5 years	NaN	Yes
4	No	I do not use machine learning methods	NaN	NaN
5	Yes	I do not use machine learning methods	MySQL	NaN
6	Yes	5-10 years	PostgreSQL	Yes
7	No	1-2 years	NaN	Yes

▼ Data Munging

.. .. I do not use machine learning

From the table above, we see that none of the columns has numeric values. In `scikit-learn`, there are limited ways of building models that directly work with categorical attributes. We need to preprocess these columns before we can build and evaluate models. More specifically, we need to encode these columns in numeric values. The 3 feature columns in this small dataset are different and we will pre-process each in a different way. In fact, they represent the three types of columns in the larger dataset. Therefore, the following tasks of pre-processing the small dataset will prepare you for working on the larger dataset.

10 yes 5-10 years Snowflake yes

▼ 1. Binary attribute: Regularly use Scikit-learn

The column `Regularly use Scikit-learn` describes whether a person uses `scikit-learn` on a regular basis. It has two values 'Yes' and NaN (i.e., null value). Based on how the dataset was created, NaN here means 'No'. Let's replace the values in this column with 1 and 0.

Task 1: In column `Regularly use Scikit-learn`, replace 'Yes' by 1. (5 points)

Hint: You can use `replace()`. There are definitely many other ways and you are free to use any approach that works. In Programming Assignment 2, there are similar tasks.

Code for Task 1

```
survey.replace(to_replace = {'Regularly use Scikit-learn': 'Yes'}, value = 1, inplace = True)
```

Task 2: In column Regularly use Scikit-learn, replace NaN by 0. (5 points)

Hint: You can use `fillna()`. There are definitely many other ways and you are free to use any approach that works. In Programming Assignment 2, there are similar tasks.

```
# Code for Task 2
```

```
survey['Regularly use Scikit-learn'].fillna(0, inplace=True)
```

If your code is correct, the results of `survey.head(20)` will be as follows. Note that the values are floating-point numbers. Feel free to use integers. It may or may not have an impact on your classifier accuracy and it is up to you to explore.

```
survey.head(20)
```

	Yearly salary	Years of experience in machine learning methods	Most frequently used big data products	Regularly use Scikit-learn
1	No	3-4 years	MySQL	1.0
3	No	4-5 years	NaN	1.0

▼ 2. Nominal attribute: Most frequently used big data products

The column `Most frequently used big data products` describes the big data product that a person uses most frequently. It has values such as `MySQL`, `PostgreSQL` and so on. Based on what we learned earlier in the semester, this is a nominal attribute in that there isn't a meaningful order among the attribute values. We will use one-hot encoding to represent this attribute. More specifically, we will make one new binary-value column for each distinct big data product. A row has value 1 or 0 in that new column, based on its value in the original `Most frequently used big data products` column. In Programm Assignment 2, we actually performed similar operations.

Go ahead to apply the following code. After that, the results of `survey.head(20)` show the new columns, each with the prefix `Bigd`. Note that we also dropped the original column `Most frequently used big data products`.

```
bigd = pd.get_dummies(survey['Most frequently used big data products'], prefix='Bigd')
survey = survey.drop(['Most frequently used big data products'], axis=1)
survey = pd.concat([survey, bigd], 1)
```

```
survey.head(20)
```

	Yearly salary	Years of experience in machine learning methods	Regularly use Scikit-learn	Bigd_Amazon Aurora	Bigd_Amazon DynamoDB	Bigd_Amazon RDS	Bigd_Amazon Redshift	
1	No	3-4 years	1.0	0	0	0	0	
3	No	4-5 years	1.0	0	0	0	0	
4	No	I do not use machine learning methods	0.0	0	0	0	0	
5	Yes	I do not use machine learning methods	0.0	0	0	0	0	
6	Yes	5-10 years	1.0	0	0	0	0	
7	No	1-2 years	1.0	0	0	0	0	
8	No	5-10 years	0.0	0	0	0	0	
9	No	Under 1 year	1.0	0	0	0	0	
10	No	I do not use machine learning methods	0.0	0	0	0	0	
11	No	4-5 years	1.0	0	0	0	0	
12	No	1-2 years	0.0	0	0	0	0	

▼ 3. Ordinal attribute: Years of experience in machine learning methods

Let's take a look at the distinct values of column Years of experience in machine learning methods . This is an ordinal attribute, since these values capture different levels of experience, from none to abundant experience. Let's map these values into the scale of 1-9 .

```
survey['Years of experience in machine learning methods'].unique()

array(['3-4 years', '4-5 years', 'I do not use machine learning methods',
      '5-10 years', '1-2 years', 'Under 1 year', '2-3 years', nan,
      '20 or more years', '10-20 years'], dtype=object)
```

Task 3: In column Years of experience in machine learning methods, replace column values by 1 - 9 --- 'I do not use machine learning methods' by numeric value 1, 'Under 1 year' by numeric value 2 and '20 or more years' by numeric value 9. (5 points)

Code for Task 3

```
survey['Years of experience in machine learning methods'] = survey['Years of experience in ma
```

The column Years of experience in machine learning methods has null values. We are going to replace these null values by 0. Note that this is not an ideal solution. Given that 0 is less than 1, the classification model we are going to build may pick up the signal that a person having 0 in this column has less experience than a person having 1, which may not be the case. However, we don't really have a better solution, unless we keep the null values. There are some implementation of learning algorithms in `scikit-learn` that admit null values and there are other libraries to use. But let's don't make things too complicated in this assignment. Let's just replace NaN by 0 in this column.

Task 4: In column Years of experience in machine learning methods, replace NaN by 0. (5 points)

Code for Task 4

```
survey['Years of experience in machine learning methods'].fillna(0, inplace=True)
```

If your code is correct, the results of `survey.head(20)` will be as follows.

```
survey.head(20)
```


	Yearly salary	Years of experience in machine learning methods	Regularly use Scikit-learn	Bigd_Amazon Aurora	Bigd_Amazon DynamoDB	Bigd_Amazon RDS	Bigd_Amazon Redshift	l
1	No	5.0	1.0	0	0	0	0	
3	No	6.0	1.0	0	0	0	0	
4	No	1.0	0.0	0	0	0	0	
5	Yes	1.0	0.0	0	0	0	0	
6	Yes	7.0	1.0	0	0	0	0	
7	No	3.0	1.0	0	0	0	0	
8	No	7.0	0.0	0	0	0	0	
9	No	2.0	1.0	0	0	0	0	
10	No	1.0	0.0	0	0	0	0	
11	No	6.0	1.0	0	0	0	0	
12	No	3.0	0.0	0	0	0	0	
13	Yes	4.0	1.0	0	0	0	0	
15	No	3.0	1.0	0	0	0	0	
17	No	4.0	1.0	0	0	0	0	
18	Yes	7.0	1.0	0	0	0	0	

▼ Prepare the Larger Dataset

Now that we have finished the exercise of pre-processing the smaller dataset, let's get the larger dataset ready. Once again, the CSV files can be also found in the assignment's entry in canvas. You will need to upload these CSV files to your Google Colab working directory. Once the CSV files are in your working directory, let's load the larger CSV file `ds_workers_learners.csv` into a pandas DataFrame.

```
survey = pd.read_csv('ds_workers_learners.csv', delimiter=',', decimal=",")
```

▼ Task 5: Pre-process the larger dataset. (20 points)

The larger dataset has much more columns than the smaller one. However, they are similar to the 3 types of columns we explained earlier. Go ahead to preprocess these columns. Furthermore, get the class attribute `Yearly salary` ready in the same way as we did on the smaller dataset.


```
bigd = pd.get_dummies(survey['Most frequently used big data products'], prefix='Bigd')
survey = survey.drop(['Most frequently used big data products'], axis=1)
survey = pd.concat([survey, bigd], 1)

tool = pd.get_dummies(survey['Primary tool for analyzing data'], prefix='Tool')
survey = survey.drop(['Primary tool for analyzing data'], axis=1)
survey = pd.concat([survey, tool], 1)

survey.replace(to_replace = {'Regularly use Python': 'Yes'}, value = 1, inplace = True)
survey['Regularly use Python'].fillna(0, inplace=True)

survey.replace(to_replace = {'Regularly use R': 'Yes'}, value = 1, inplace = True)
survey['Regularly use R'].fillna(0, inplace=True)

survey.replace(to_replace = {'Regularly use SQL': 'Yes'}, value = 1, inplace = True)
survey['Regularly use SQL'].fillna(0, inplace=True)

survey.replace(to_replace = {'Regularly use Scikit-learn': 'Yes'}, value = 1, inplace = True)
survey['Regularly use Scikit-learn'].fillna(0, inplace=True)

survey.replace(to_replace = {'Regularly use TensorFlow': 'Yes'}, value = 1, inplace = True)
survey['Regularly use TensorFlow'].fillna(0, inplace=True)

survey.replace(to_replace = {'Regularly use Keras': 'Yes'}, value = 1, inplace = True)
survey['Regularly use Keras'].fillna(0, inplace=True)

survey.replace(to_replace = {'Regularly use PyTorch': 'Yes'}, value = 1, inplace = True)
survey['Regularly use PyTorch'].fillna(0, inplace=True)

survey.replace(to_replace = {'Regularly use Xgboost': 'Yes'}, value = 1, inplace = True)
survey['Regularly use Xgboost'].fillna(0, inplace=True)

survey.replace(to_replace = {'Regularly use Linear or Logistic Regression': 'Yes'}, value = 1, inplace = True)
survey['Regularly use Linear or Logistic Regression'].fillna(0, inplace=True)

survey.replace(to_replace = {'Regularly use Decision Trees or Random Forests': 'Yes'}, value = 1, inplace = True)
survey['Regularly use Decision Trees or Random Forests'].fillna(0, inplace=True)

survey.replace(to_replace = {'Regularly use Gradient Boosting Machines': 'Yes'}, value = 1, inplace = True)
survey['Regularly use Gradient Boosting Machines'].fillna(0, inplace=True)

survey.replace(to_replace = {'Regularly use Bayesian Approaches': 'Yes'}, value = 1, inplace = True)
survey['Regularly use Bayesian Approaches'].fillna(0, inplace=True)

survey.replace(to_replace = {'Regularly use Convolutional Neural Networks': 'Yes'}, value = 1, inplace = True)
survey['Regularly use Convolutional Neural Networks'].fillna(0, inplace=True)
```

With our solution, the first 20 rows after pre-processing the dataset will look like the following. Note that you don't necessarily need to handle each column in exactly the same way as the teaching staff do. In certain cases there could be multiple sensible choices.

```
survey.head(20)
```

	Age	Degree	Size of employer	Yearly salary	Years of coding experience	Years of experience in machine learning methods	Experience with TPU	Regularly use Python	Regul u
1	8.0	4	1	No	5	5.0	1.0	1.0	
3	3.0	4	5	No	4	6.0	1.0	1.0	
4	5.0	3	4	No	1	1.0	1.0	0.0	
5	7.0	4	5	Yes	6	1.0	4.0	1.0	
6	4.0	6	5	Yes	5	7.0	3.0	1.0	
7	2.0	4	4	No	2	3.0	1.0	1.0	
8	8.0	3	2	No	6	7.0	2.0	1.0	
9	3.0	4	5	No	2	2.0	1.0	1.0	
10	8.0	3	4	No	2	1.0	1.0	1.0	
11	5.0	4	2	No	4	6.0	3.0	1.0	
12	8.0	4	4	No	6	3.0	1.0	1.0	
13	10.0	4	1	Yes	3	4.0	1.0	1.0	
15	10.0	4	4	No	6	3.0	1.0	1.0	
17	9.0	4	1	No	6	4.0	1.0	1.0	
18	3.0	4	1	Yes	5	7.0	3.0	1.0	
19	3.0	4	4	No	0	0.0	0.0	0.0	
20	8.0	4	5	Yes	6	7.0	1.0	1.0	
21	3.0	2	5	No	1	2.0	2.0	0.0	
26	1.0	4	3	No	3	1.0	1.0	1.0	
28	9.0	4	2	No	6	4.0	1.0	1.0	

20 rows × 93 columns

▼ Load Pre-processed Dataset

If you couldn't get Task 5 done, don't panic. We provide a preprocessed file `p3_processed.csv` to you, which is in the same place in Canvas as the small/large dataset files. You just need to run the following code to load it. In fact, you should use this preprocessed data file regardless, even if you successfully finish your Task 5. This way we make sure everyone uses the same data file for creating the classification models, which allows us to fairly grade all submissions.

```
survey = pd.read_csv('p3_processed.csv', delimiter=',')
```

▼ Task 6: Build and evaluate classification models. (45 points)

You can apply any of the methods explained in the [instruction notebook](#) or any other method in `scikit-learn`. You can choose which feature columns to include in building the model. You can tune your model by using any combination of parameter values. You can even implement your own method.

Make sure to follow the good practice we learned about model selection and model evaluation. For model evaluation:

1. Partition the dataset into training set and test set. The test set shouldn't be used in any way during training your model.
2. Use cross-validation in order to get more robust evaluation results.
3. After evaluation, you can train your model again on the whole dataset. Then the trained model can be made available to classify unseen instances in the future. Of course, in this assignment, we don't really have unseen instances to be applied. Maybe you can plug in your own information to see how the model predicts, just for fun.

For model selection:

1. Model selection is the step for choosing the optimal model among multiple different types of models (e.g., a decision tree vs. a kNN classifier), or for tuning the hyperparameters (e.g., the maximum depth in a decision tree) in order to get the optimal model within the same family of models.
2. In model selection, you further partition the training set (from model evaluation) into train set and validation set. (Here we call it 'train set', to make it clear it is a subset of the 'training set'.)
3. Different models are trained using the train test and their performance on the validation set is used to select the best model and/or best hyperparameters.
4. Model section itself can also use cross-validation.

Note that it is non-trivial to implement model selection on your own. Fortunately, `scikit-learn` provides support for this too. **In this assignment, you are required to perform model evaluation. You are not required to perform model selection.** Instead, in this assignment, you can compare and select models based on their performance on the test set. (When deploying production model, this leads to overfitting and thus should be avoided. But it is fine for this assignment.)

In the code block below, you will find the baseline results from the teaching staff's code.

Code for Task 6

```
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import OrdinalEncoder
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import LinearSVC
from sklearn.naive_bayes import GaussianNB
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import classification_report

survey = survey.loc[:, ~survey.columns.str.startswith('Gender')]
survey = survey.loc[:, ~survey.columns.str.startswith('Industry')]

class_column = 'Yearly salary'

feature_columns = survey.columns.values
feature_columns = feature_columns[feature_columns != 'Yearly salary']

survey_feature = survey[feature_columns]
survey_class = survey[class_column]

print("First three rows")
print(survey[0:3])
print()
print("Feature columns, first three rows:")
print(survey_feature[0:3])
print()
print("Class column, first three rows:")
print(survey_class[0:3])
print()
```



First three rows

Age	Degree	...	Tool_Local development environments	Tool_Other
0	8	4	...	0
				0

1	3	4	...	0	0
2	5	3	...	0	0

[3 rows x 70 columns]

Feature columns, first three rows:

	Age	Degree	...	Tool_Local development environments	Tool_Other
0	8	4	...	0	0
1	3	4	...	0	0
2	5	3	...	0	0

[3 rows x 69 columns]

Class column, first three rows:

0	No
1	No
2	No

Name: Yearly salary, dtype: object

```
train_feature, test_feature, train_class, test_class = \
    train_test_split(survey_feature, survey_class, stratify=survey_class, \
        train_size=0.75, test_size=0.25)
```

```
#model = DecisionTreeClassifier()
#model = KNeighborsClassifier(n_neighbors=11)
#model = LinearSVC(random_state=0, tol=1e-1, max_iter=1000)
model = MultinomialNB()
```

```
scores = cross_val_score(model, survey_feature, survey_class, cv=5)
print("Cross-validation scores: {}".format(scores))
print("Average cross-validation score: {:.2f}".format(scores.mean()))
print()
model.fit(train_feature, train_class)
print("Test set accuracy: {:.2f}".format(model.score(test_feature, test_class)))
print()
prediction = model.predict(test_feature)
#print("Test set predictions:\n{}".format(prediction))
print()
print("Confusion matrix:")
print(pd.crosstab(test_class, prediction, rownames=['True'], colnames=['Predicted'], margins=
print()
print("Classification report:")
print(classification_report(test_class, prediction))
```

```
Cross-validation scores: [0.75880759 0.75338753 0.7696477  0.74254743 0.76151762]
Average cross-validation score: 0.76
```

```
Test set accuracy: 0.75
```

Confusion matrix:

Predicted	No	Yes	All
-----------	----	-----	-----

```

True
No      269   70  339
Yes     45   78  123
All     314  148  462

```

Classification report:

	precision	recall	f1-score	support
No	0.86	0.79	0.82	339
Yes	0.53	0.63	0.58	123
accuracy			0.75	462
macro avg	0.69	0.71	0.70	462
weighted avg	0.77	0.75	0.76	462

▼ Explain Your Work

Task 7: Document and explain your models and results. (15 points)

You are required to write a **brief document** (500-1000 words, not including the words in tables and figures of evaluation results) to discuss the process you went through to explore and compare different methods and choose parameter values. The document should also report the classification accuracy evaluation results of different methods you have experimented with and discuss how you finalize your choice. Make sure to include performance measures such as confusion matrix, classification report, cross validation scores, and test set accuracy that are returned from the code. You are encouraged to include other presentations of evaluation results.

Write your document below. Note that you may need to use multiple code and text blocks in order to produce a document with evaluation results.

Grading Rubrics

Your tasks 6 and 7 will be evaluated on correctness, classification accuracy, efficiency, report quality, and code quality. Make sure to thoroughly understand the following grading rubrics.

(1) **Basics:** 10 points

You will be evaluated on whether you can accomplish the given tasks, i.e., a complete classification model.

(2) **Execution efficiency:** 10 points

10 points: your code finishes in seconds (this can vary for different methods and we will take that into consideration).

5 points: your code is clearly much slower than majority of the submissions.

0 points:* your code will need to take hours to finish, OR mostly incorrect implementation which makes efficiency evaluation not meaningful.

(3) Accuracy: 15 points

How much can you improve your classifier's accuracy to outperform our baseline, which is some very primitive model without any effort of improvement. We will run the codes of all students and compare your classifiers' performance.

15 points: Among the best performance in the class AND performance clearly better than the baseline results.

**12 points:* Stronger performance than majority of the class AND performance slightly stronger than the baseline results.

9 points: Average performance in the class AND performance on par with the baseline results.

6 points: weaker performance than majority of the class OR performance weaker than the baseline results.

3 points: Weaker performance than 85% of the class OR performance clearly weaker than the baseline results.

0 points: Mostly incorrect implementation which makes accuracy evaluation not meaningful.

(4) Report---clarity, organization, correctness, thoroughness: 15 points

Your report will be graded based on whether it is clear and well organized, whether the discussion is correct and logical, whether it demonstrates proper understanding of classification, and whether it is thorough in presenting designing choices and results.

(5) Quality---clarity, organization, modularity, comments: 10 points

Follow good coding standards to make your program easy to understand and easy to maintain/extend. Provide sufficient comments in your code and make it self-explaining.

High mark : 10 points

Medium mark : 7 points

Low mark : 3 points

Poor: 0 points

(4) Total score of Tasks 6 and 7: 60 points

Your score will be calculated from the individual break-ups using the following equation:

Basics + Efficiency + Accuracy +Report + Quality

